

Probabilistic Linear Solvers for Machine Learning

Jonathan Wenger

Philipp Hennig

NeurIPS 2020

EBERHARD KARLS
UNIVERSITÄT
TÜBINGEN




Max Planck Institute for
Intelligent Systems
imprs-is



erc

some of the presented work is supported
by the European Research Council.






$$\mathbf{A} \mathbf{x}_* = \mathbf{b}$$


where $\mathbf{A} \in \mathbb{R}^{n \times n}$ symmetric positive definite.

ML-specific Challenges


- large-scale systems
- characteristic structure
- generative information
- subject to noise




general linear model $\mathbf{X}^\top \mathbf{X}$



kernel matrix $\mathbf{K} = k_{\text{RBF}}(\mathbf{X}, \mathbf{X})$



Hessian matrix $\mathbf{H} = \nabla^2 f(\boldsymbol{\theta})$



Algorithm Components

1. Prior $p(\mathbf{A})$ or $p(\mathbf{H})$
2. Policy $\mathbf{s}_i = -\mathbb{E}[\mathbf{H}](\mathbf{A}\mathbf{x}_{i-1} - \mathbf{b})$
3. Observation $\mathbf{y}_i = \mathbf{A}\mathbf{s}_i$
4. Belief Update for \mathbf{A} and \mathbf{H}
5. (Uncertainty Calibration)

Theoretical Properties

- ✦ Conjugate Directions Method ($\leq n$ iters)
- ✦ Recovers CG given certain priors
- ✦ Time complexity $\mathcal{O}(kn^2)$
- ✦ Space complexity $\mathcal{O}(kn)$

Desiderata

No.	Property	
(1)	distribution over matrices	✓
(2)	symmetry	✓
(3)	positive definiteness	~
(4)	positive linear combination in same distr. family	✓
(5)	corresponding priors on matrix and inverse	✓
⋮	⋮	⋮

Covariance Class

$$p(\mathbf{A}) = \mathcal{N}(\mathbf{A}; \mathbf{A}_0, \mathbf{W}_0^{\mathbf{A}} \otimes \mathbf{W}_0^{\mathbf{A}})$$


$$p(\mathbf{H}) = \mathcal{N}(\mathbf{H}; \mathbf{A}_0^{-1}, \mathbf{W}_0^{\mathbf{H}} \otimes \mathbf{W}_0^{\mathbf{H}})$$

where $\mathbf{W}_0^{\mathbf{A}}$ and $\mathbf{W}_0^{\mathbf{H}}$ admit degrees of freedom for *uncertainty calibration*.

Prior Spectral Knowledge

Spectrum $\lambda(\mathbf{A})$

$$\text{Rayleigh Quotient } R(\mathbf{A}, \mathbf{s}) = \frac{\mathbf{s}^\top \mathbf{A} \mathbf{s}}{\mathbf{s}^\top \mathbf{s}}$$






ProbNum implements probabilistic numerical methods in Python.


<https://github.com/probabilistic-numerics/probnum>

or alternatively `pip install probnum`.

Future Applications



Galerkin Methods $Au = f$




Empirical Risk Minimization $Hd = g$



Probabilistic Linear Solvers for Machine Learning

Jonathan Wenger and Philipp Hennig

- ✦ Linear systems in ML exhibit characteristic structure.
- ✦ Probabilistic linear solvers make use of prior generative information.
- ✦ Limited computational resources induce *numerical uncertainty*.



Preprint

arXiv

2010.09691

Implementation



<https://github.com/probabilistic-numerics/probnum>

Experiments



<https://github.com/JonathanWenger/probabilistic-linear-solvers-for-ml>